

# FEDERATING CONSISTENCY FOR PARTITION PRONE NETWORKS


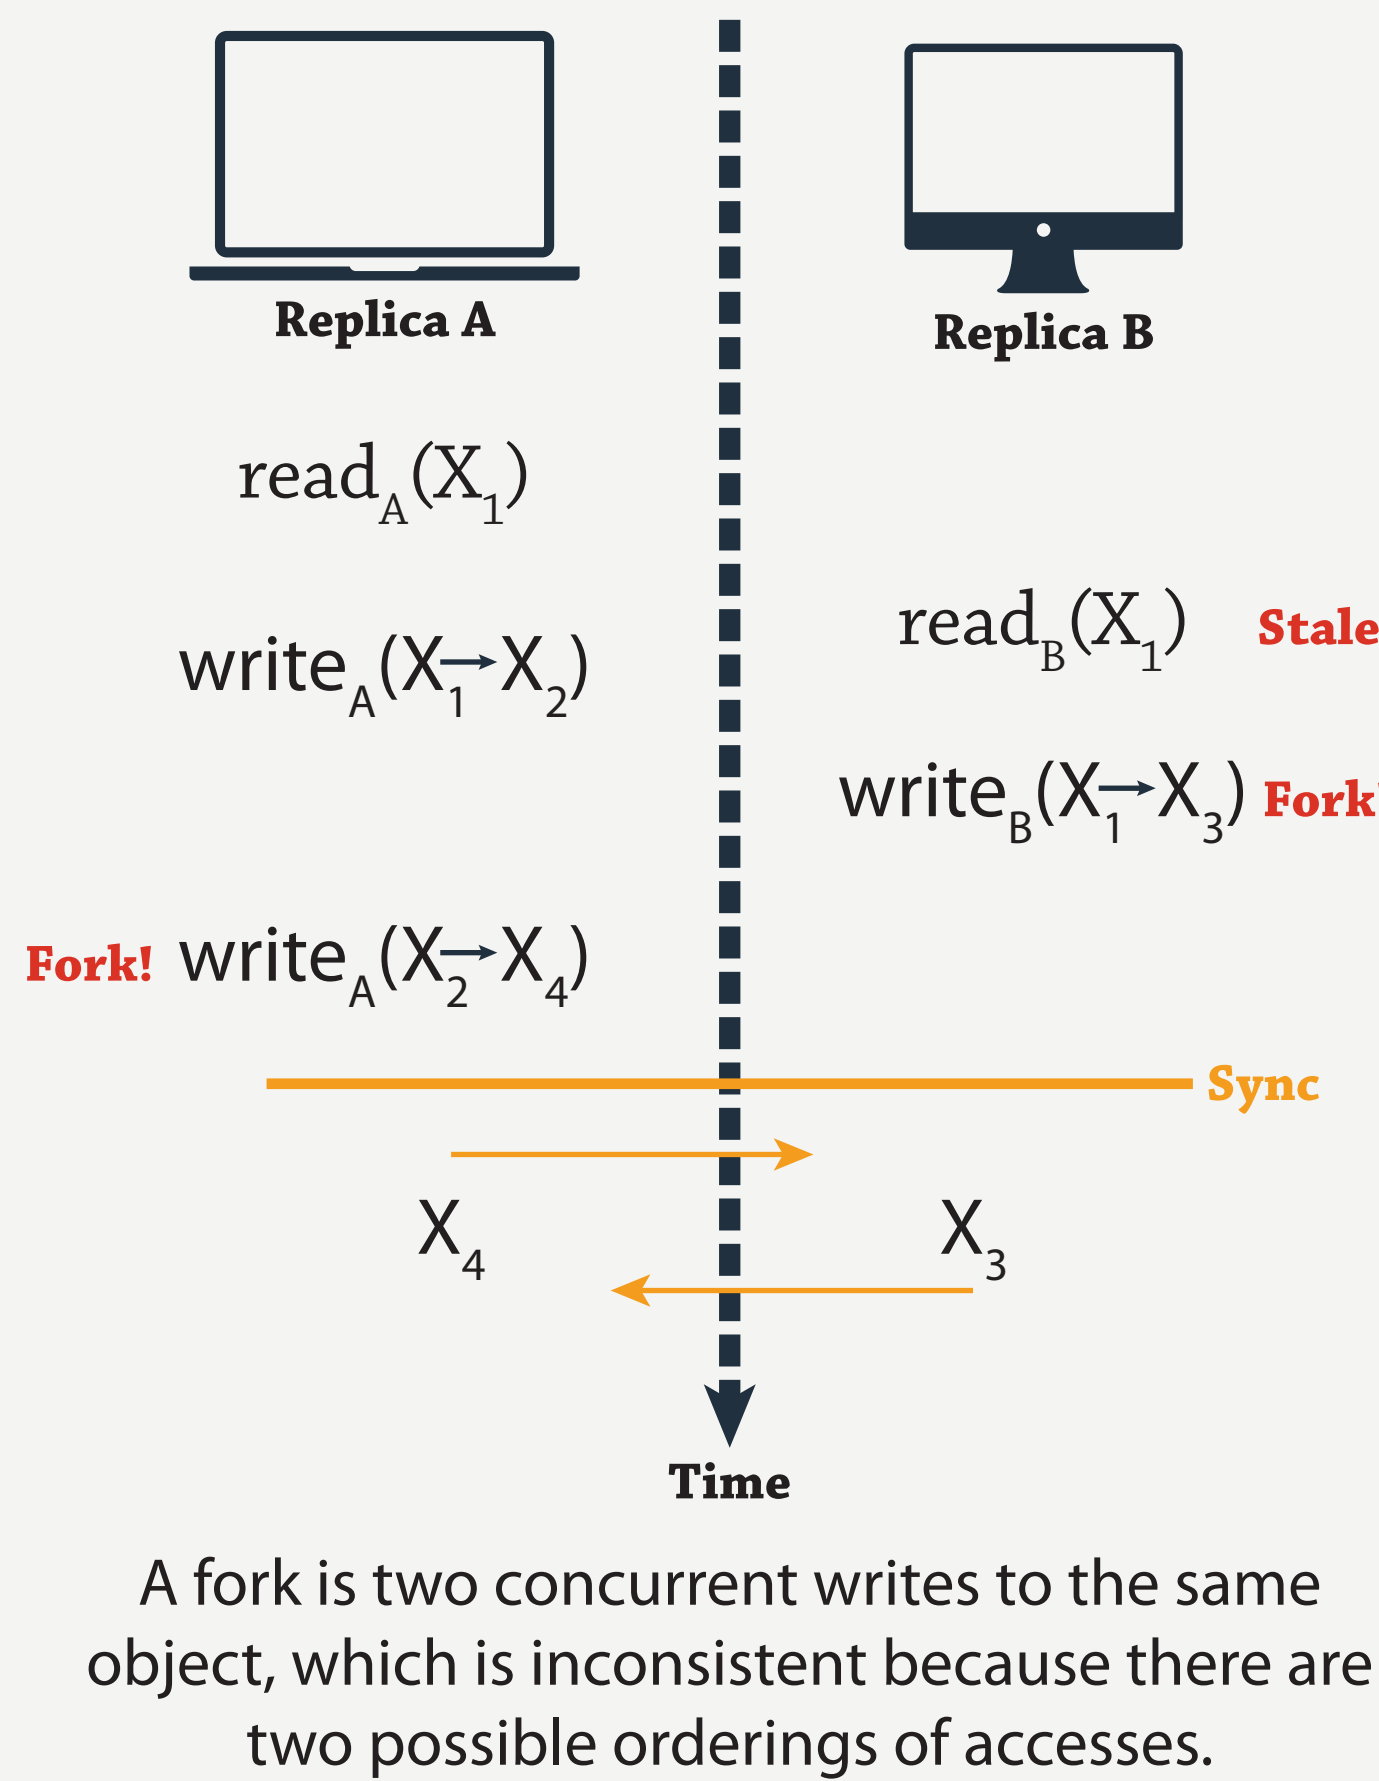
Benjamin Bengfort and Pete Keleher  
University of Maryland • Department of Computer Science



## Introduction

Groups of devices coordinated by **consensus** can efficiently order events under ideal conditions, but become less effective in dynamic and heterogenous environments. Weakly consistent devices **efficiently tolerate faults and dynamic conditions** but are slow to converge on a single ordering of system events.

Federated consistency combines the strengths of both approaches into a single protocol. Federated systems use a **strongly consistent inner core** to maintain a totally ordered sequence of accesses. A **cloud of weakly-consistent devices** disseminates orderings and enables progress despite varying connectivity.




We target dynamic, geo-replicated distributed system models that co-locate replicas with clients rather than traditional cloud-service oriented approaches. Our model considers **extremely variable latency**, bandwidth, and availability across wide-area links that connect more stable local networks.

We posit a **file system** as the natural use case of local, client-oriented systems, though the model easily generalizes to any distributed storage system. Clients access objects locally, creating a **metadata version history** that is replicated to the entire system, though data may only be partially replicated.

## Raft Strong Consistency

Leader-oriented consensus algorithms nominate a dedicated proposer in a fault-tolerant manner to conduct consistent replication. All accesses go through the leader so linearizability or sequential consistency is guaranteed. On a conflict, the leader simply drops the conflicting version.



POLICY

First Write to Leader Wins

## Hybridization of Consistency by Replica Protocol

**Consistency Integration**  
A vector component called the forte is only incremented by the Raft leader on commit.

Eventual replicas must propagate forte bumps to all derived versions.



Adjust topology and pairwise selection likelihoods

{SCALAR.PID.FORTE}  
CONFLICT FREE VERSION  
+  
FORTE COMPONENT

Federated Tunable Consistency

**Protocol Integration**  
Replicas select their local consistency policy but all Raft replicas must also participate in anti-entropy.


Eventual replicas can select frequency of Raft synchronization.



Adjust timing parameters and adapt to environment

## Anti-Entropy Eventual Consistency


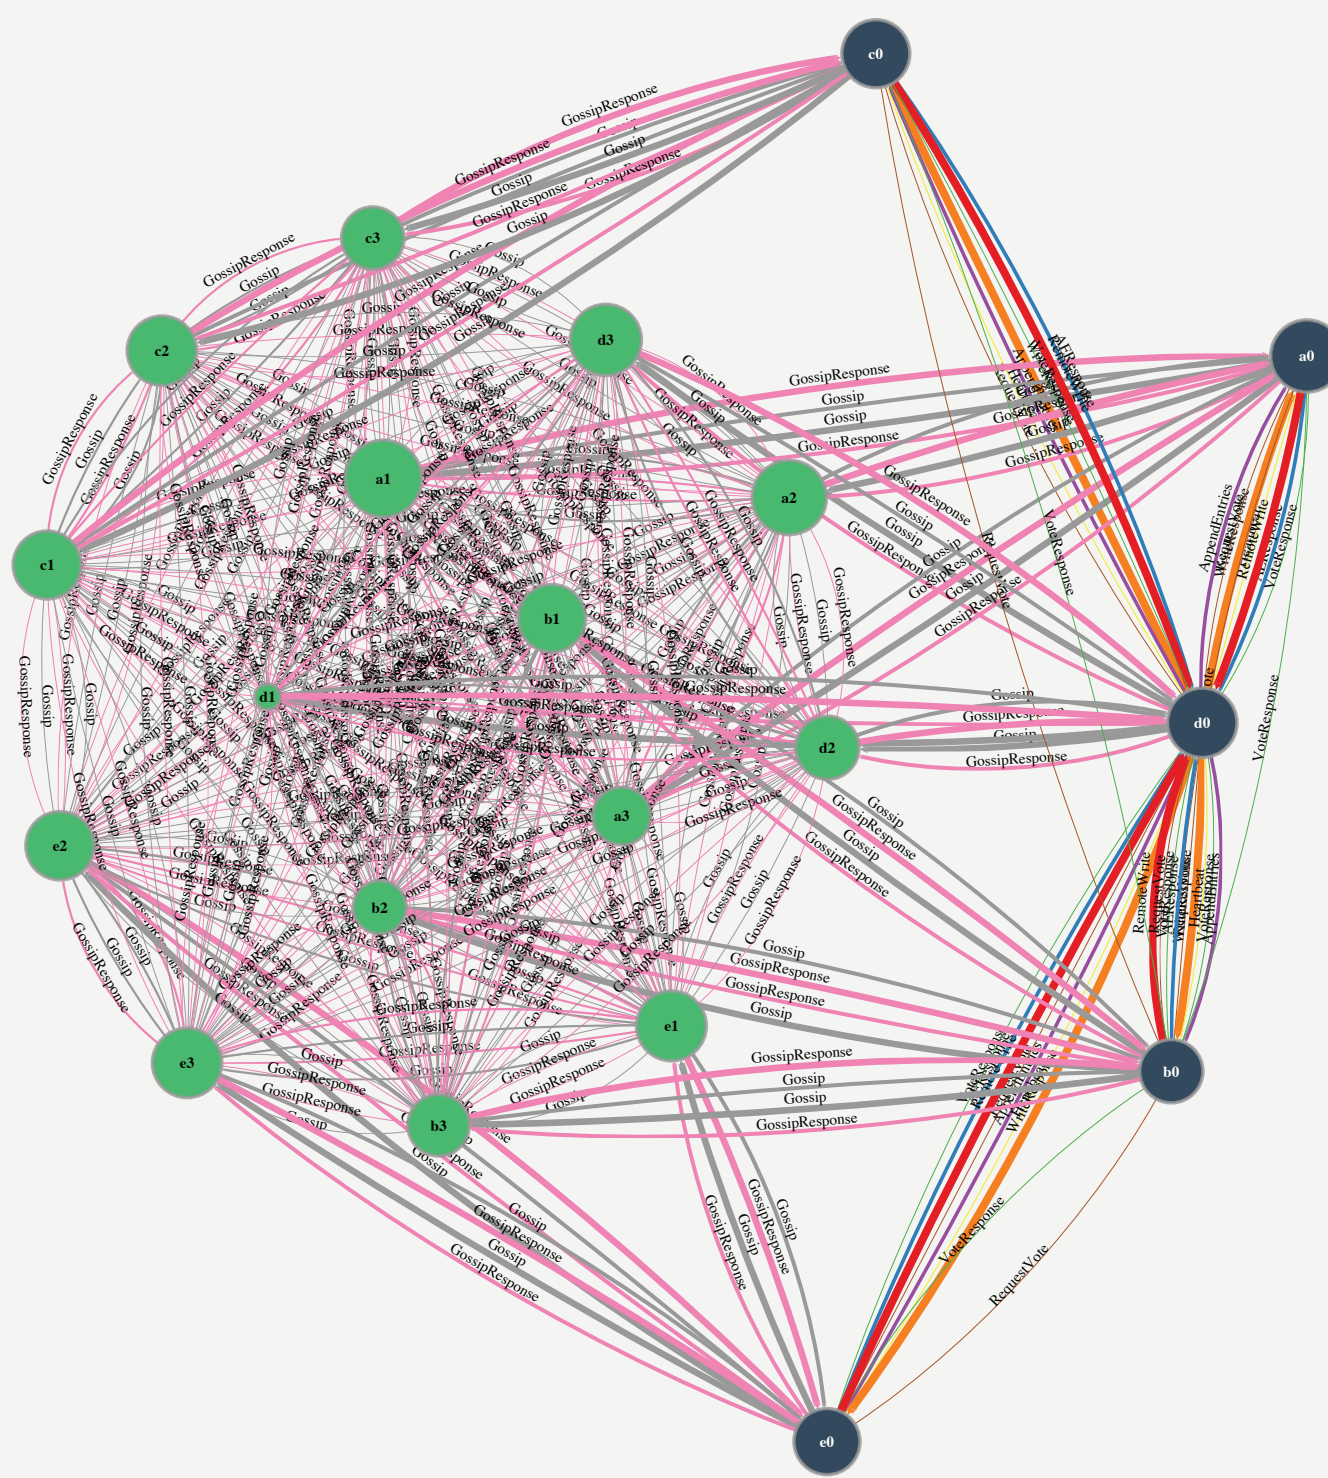
Eventually consistent replication protocols utilize periodic pairwise anti-entropy sessions to synchronize replica states. Because there is no central authority, accesses are highly available and localized. However, anti-entropy tolerates inconsistencies and may even "stomp" writes by not fully replicating them.



POLICY

Latest Version Wins

## Simulation



Our initial evaluation utilized discrete event simulations with asynchronous processes to investigate the **effect of latency, conflict, and network outages on consistency**. Every experiment includes multiple simulated runs of 20 processes in varying environments and configurations, accessing multiple objects concurrently.

The primary simulation parameter was variable **latency across wide areas**, described by the mean and standard deviation of one-way message latency. Other parameters include the **likelihood of conflict & outages**, timing parameters for consistency protocols & workload, and distributions for anti-entropy selection.

## Results

